

Inverse probability of censoring weighting (IPCW) for linear regression

Brice Ozenne

August 4, 2021

1 Principle

Inverse probability of censoring weighting (IPCW) is a method able to handle informative drop-out. Intuitively, in presence of informative drop-out a complete case analysis is a biased approach as individuals with complete data are not representative of the population. However with an appropriate re-weighting of the individuals with complete data, we can "re-balance" our sample and make it representative of the population. To do so, we divide the population into sub-populations and attribute weights to individuals who did not drop-out inversely proportional to the frequency of the drop-out in the sub-population. Thanks to the weights, individuals who did not drop-out "represent" the individuals who dropped-out. Thus, overall, the weighted sample is representative of the population.

2 Continuous outcome

2.1 Illustrative example

Consider a study where we follow depressed individual over time. They have a baseline measurement, then are given a treatment, and then have a follow-up measurement. We would like to assess the treatment effect in term of depression score ¹. The population of interest contain severely and moderately depressed individuals; the treatment may work differently in each sub-population. Unfortunately, some study participants dropped-out and it seems that they are more likely to drop-out when they are severely depressed.

¹:To simplify, there is no control group - we assume that without treatment the depression score would be constant.

We can simulate such a dataset using the following function:

```
simTrial <- function(n, rho, dmU, pC){
  require(mvtnorm)
  require(data.table)
  ## simulate data
  Sigma <- 10^2*matrix(c(1,rho,rho,1),2,2)
  ## gather into dataset
  M.Ym <- rmvnorm(n, mean = c(50, 50-dmU[1]), sigma = Sigma)
  M.Ys <- rmvnorm(n, mean = c(75, 75-dmU[2]), sigma = Sigma)
  dtL <- rbind(
    data.table(id = 1:n, mdd = "moderate", time = "T1", Y = M.Ym[,1]),
    data.table(id = 1:n, mdd = "moderate", time = "T2", Y = M.Ym[,2]),
    data.table(id = n+(1:n), mdd = "severe", time = "T1", Y = M.Ys[,1]),
    data.table(id = n+(1:n), mdd = "severe", time = "T2", Y = M.Ys[,2])
  )
  dtL$probaD0 <- 0
  dtL[time=="T2", probaD0 := ifelse(.SD$mdd=="moderate",pC[1],pC[2])]
  dtL[,dropout := rbinom(.N,prob=probaD0,size=1)]
  dtL[,Yobs:=Y]
  dtL[dropout==1,Yobs:=NA]
  dtL$probaD0 <- NULL
  return(dtL)
}
set.seed(11)
dtL <- simTrial(n = 1000, rho = 0.8, dmU = c(25,50), pC = c(0.2,0.7))
print(dtL)
```

	id	mdd	time	Y	dropout	Yobs
1:	1	moderate	T1	44.83259	0	44.83259
2:	2	moderate	T1	30.34157	0	30.34157
3:	3	moderate	T1	56.36308	0	56.36308
4:	4	moderate	T1	64.63341	0	64.63341
5:	5	moderate	T1	45.10048	0	45.10048

3996:	1996	severe	T2	30.59793	1	NA
3997:	1997	severe	T2	18.97725	1	NA
3998:	1998	severe	T2	29.80266	1	NA
3999:	1999	severe	T2	30.26518	0	30.26518
4000:	2000	severe	T2	39.15797	0	39.15797

Here we have simulated a two sub-populations of 1000, with a correlation of 0.5 between baseline and follow-up . The treatment effect is twice bigger for the severely depressed population but individuals from this population are also much more likely to drop-out. Overall the expected treatment effect is:

```
(-25-50)/2
```

```
[1] -37.5
```

Without drop-out, we could use a simple linear model to carry-out the analysis:

```
dtW.oracle <- dcast(dtL, formula = id ~ time, value.var = "Y")
dtW.oracle$diff <- dtW.oracle$T2-dtW.oracle$T1
e.oracle <- lm(diff~1, data = dtW.oracle)
summary(e.oracle)$coef
```

```
              Estimate Std. Error   t value Pr(>|t|)
(Intercept) -37.35098   0.3141814 -118.8835      0
```

leading to an estimate quite close to the true value.

With drop-out, a complete case analysis would lead to a biased estimator. In this example, we can "see" that the estimated value is far away from the true one (even when accounting for the uncertainty):

```
dtW <- dcast(dtL, formula = id + mdd ~ time, value.var = "Yobs")
dtW$diff <- dtW$T2-dtW$T1
dtW.CC <- dtW[!is.na(diff)]
e.CC <- lm(diff~1, data = dtW.CC)
summary(e.CC)$coef
```

```
              Estimate Std. Error   t value Pr(>|t|)
(Intercept) -31.42356   0.3909029 -80.38713      0
```

An alternative approach would be to use a linear mixed model (i.e. full information):

```
require(nlme)
e.FI <- lme(Yobs~time, random = ~1|id, data = dtL, na.action = na.omit)
summary(e.FI)$tTable
```

```
              Value Std.Error   DF   t-value p-value
(Intercept)  62.59128 0.3239587 1999 193.20760      0
timeT2       -33.76472 0.3855964 1068 -87.56494      0
```

which is better than the complete case analysis still biased when the drop-out mechanism depends on variables other than the baseline value.

A better approach is to use IPCW. First we model the probability of not dropping out at follow-up:

```
dtW$observed <- !is.na(dtW$T2)
e.glmW.oracle <- glm(observed ~ mdd, data = dtW,
  family = binomial(link = "logit"))
```

and then compute the weights for observations with full data:

```
dtW$weight.oracle <- 1/predict(e.glmW.oracle, newdata = dtW, type = "
  response")
dtW[observed == TRUE, sum(weight.oracle)]
```

```
[1] 2000
```

Note that the weights sum to the total sample size. We then perform the complete case analysis with these weights:

```
dtW.CC <- dtW[!is.na(diff)]
e.IPCW.oracle <- lm(diff~1, data = dtW.CC, weights = dtW.CC$weight.oracle)
summary(e.IPCW.oracle)$coef
```

```
              Estimate Std. Error  t value Pr(>|t|)
(Intercept) -36.89889   0.4251421 -86.7919      0
```

which gives a result much closer to the true value. A more feasible IPCW would use the baseline score to define the weights:

```
e.glmW <- glm(observed ~ T1, data = dtW,
  family = binomial(link = "logit"))
dtW$weight <- 1/predict(e.glmW, newdata = dtW, type = "response")
dtW[observed == TRUE, sum(weight)]
```

```
[1] 2015.739
```

We then perform the complete case analysis with these new weights:

```
dtW.CC <- dtW[!is.na(diff)]
e.IPCW <- lm(diff~1, data = dtW.CC, weights = dtW.CC$weight)
summary(e.IPCW)$coef
```

```
              Estimate Std. Error  t value Pr(>|t|)
(Intercept) -35.47206   0.423423 -83.77453      0
```

2.2 Simulation study

The quality of the previous estimators is compared using a simulation study. The results are summarized by [Figure 1](#).




Figure 1: Comparison between the empirical distributions of the estimators (Gaussian case) for a sample size of 1000 using 100 datasets.

3 Binary outcome

3.1 Illustrative example

A somehow similar approach can be used for binary endpoints. Consider now a study comparing the survival probability at 1 year of patients treated with a new drug vs. standard care. The population is composed of two types of patients, say some with hypertension and some without. Survival as well as the treatment effect may differ depending of the hypertension status. Hypertension may also affect the drop-out probability.

We can simulate such a dataset using the following function:

```
simTrial <- function(n, dmu, dpC){
  require(BuyseTest)
  require(data.table)
  ## simulate data
  dt1 <- simBuyseTest(n.T = n, n.C = n,
    argsBin = NULL, argsCont = NULL,
    argsTTE = list(scale.T = 1+dmu[1],
      scale.C = 1,
      scale.Censoring.T = 1+dpC[1],
      scale.Censoring.C = 1),
    latent = TRUE)
  dt2 <- simBuyseTest(n.T = n, n.C = n,
    argsBin = NULL, argsCont = NULL,
    argsTTE = list(scale.T = 2+dmu[2],
      scale.C = 2,
      scale.Censoring.T = 2+dpC[2],
      scale.Censoring.C = 2),
    latent = TRUE)
  ## gather into dataset
  dt <- rbind(
    cbind(id = 1:NROW(dt1), group = "G1", dt1),
    cbind(id = NROW(dt1) + 1:NROW(dt2), group = "G2", dt2)
  )
  return(dt)
}
set.seed(11)
tau <- 1

dt <- simTrial(n = 10000, dmu = c(0,1), dpC = c(0,1))
dt$responseUncensored <- dt$eventtimeUncensored<=tau
dt$response <- ifelse((dt$status==1)+(dt$eventtime>tau),dt$eventtime<=tau,
  NA)
dt$observed <- ifelse((dt$status==1)+(dt$eventtime>tau),1,0)
print(dt)
```

	id	group	treatment	eventtimeUncensored	eventtimeCensoring
1:	1	G1	C	1.63238841	1.74283865
2:	2	G1	C	0.08938341	1.10407172
3:	3	G1	C	1.54194414	1.11212966
4:	4	G1	C	1.04592013	1.00584279
5:	5	G1	C	0.55276522	0.04955419

39996:	39996	G2	T	4.08919433	1.23091105

39997:	39997	G2	T	3.59736307	8.14939225
39998:	39998	G2	T	7.21110232	3.04114191
39999:	39999	G2	T	0.06096057	0.42120185
40000:	40000	G2	T	6.89615584	2.22637070

	eventtime	status	responseUncensored	response	observed
1:	1.63238841	1	FALSE	FALSE	1
2:	0.08938341	1	TRUE	TRUE	1
3:	1.11212966	0	FALSE	FALSE	1
4:	1.00584279	0	FALSE	FALSE	1
5:	0.04955419	0	TRUE	NA	0

39996:	1.23091105	0	FALSE	FALSE	1
39997:	3.59736307	1	FALSE	FALSE	1
39998:	3.04114191	0	FALSE	FALSE	1
39999:	0.06096057	1	TRUE	TRUE	1
40000:	2.22637070	0	FALSE	FALSE	1

In absence of drop-out, we can compare the survival probabilities at 1 year using a logistic regression:

```
e.oracle <- glm(responseUncensored ~ treatment,
  data = dt, family = binomial(link="logit"))
summary(e.oracle)$coef
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.07022885	0.01415085	4.96287	6.945906e-07
treatmentT	-0.23721580	0.02004104	-11.83650	2.527721e-32

In presence of (differential) drop-out, a complete case analysis (i.e. restricting the analysis to the patients where the survival status at 1 year is known) would be biased:

```
dt.cc <- dt[dt$observed==1]
e.cc <- glm(response ~ treatment,
  data = dt.cc, family = binomial(link="logit"))
summary(e.cc)$coef
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.4023018	0.01814959	22.16589	7.330653e-109
treatmentT	-0.3656280	0.02510139	-14.56605	4.618541e-48

A first idea would be to re-use the IPCW approach, first fitting a logistic model for the probability of being observed at 1-year and then computing the weights:

```
e.IPCmodel <- glm(observed ~ group*treatment, data = dt, family = binomial
  (link="logit"))
dt$IPCweights <- 1/predict(e.IPCmodel, newdata = dt, type = "response")
sum(dt$IPCweights)
```

```
[1] 62570.28
```

The subsequent estimator will not be correct:

```
dt.cc <- dt[dt$observed==1]
e.IPCWcc <- glm(response ~ treatment, data = dt.cc,
  family = binomial(link="logit"), weights = dt.cc$IPCweights)
summary(e.IPCWcc)$coef
```

Advarselsbesked:

```
I eval(family$initialize) : non-integer #successes in a binomial glm!
              Estimate Std. Error   z value    Pr(>|z|)
(Intercept)  0.4593548 0.01451679   31.64300 9.465106e-220
treatmentT  -0.2997806 0.02029810  -14.76889 2.324953e-49
```

as we disregarded the duration of observation among the censored individuals. Intuitively, individuals censored early are more at risk of dying and therefore should "transfer" more weight than those censored late, e.g. just before 1 year, who don't really need to transfer weights. This can be performed using a survival model (here a Cox model) and using as weights the inverse of the probability of not being censored at the earliest between when the event occurred and 1 year:

```
library(survival)
library(riskRegression)
e.IPCmodel2 <- coxph(Surv(eventtime,status==0) ~ group*treatment,
  data = dt, x = TRUE, y = TRUE)
iPred <- predictCox(e.IPCmodel2, newdata = dt,
  time = pmin(dt$eventtime,tau)-(1e-12), diag = TRUE)$survival
dt$IPCweights2 <- dt$observed/iPred
sum(dt$IPCweights2)
```

```
[1] 40028.88
```

Note that now the sum of the weights is nearly equal to the full sample. We can then use the weights in a logistic model:

```
dt.cc <- dt[dt$observed==1]
e.IPCWcc <- glm(response ~ treatment, data = dt.cc, family = binomial(link
  ="logit"), weights = dt.cc$IPCweights2)
summary(e.IPCWcc)$coef
```


Advarselsbesked:

```
I eval(family$initialize) : non-integer #successes in a binomial glm!
              Estimate Std. Error    z value    Pr(>|z|)
(Intercept)  0.06439848 0.01414818   4.551716 5.321009e-06
treatmentT  -0.23858152 0.02003591 -11.907693 1.079216e-32
```

which is very close to the true value. Note that this estimator is implemented in the riskRegression package:

```
e.wglm <- wglm(regressor.event = ~treatment,
               formula.censor = Surv(eventtime,status==0)~group*treatment,
               times = 1,
               data = dt[,.(eventtime,status,group,treatment)])
summary(e.wglm)
```

IPCW logistic regression :

```
-----
- time: 1
glm(XX_status.1_XX ~ treatment, family = binomial(link = "logit"),
    weights = "XX_IPCW.1_XX")
```

```
              Estimate Std. Error    z value    Pr(>|z|)
(Intercept)  0.06439847 0.01931759   3.333670 0.0008570818
treatmentT  -0.23858152 0.02652596  -8.994264 0.0000000000
-----
```

3.2 Simulation study

The quality of the previous estimators is compared using a simulation study. The results are summarized by [Figure 2](#) and [Figure 3](#)

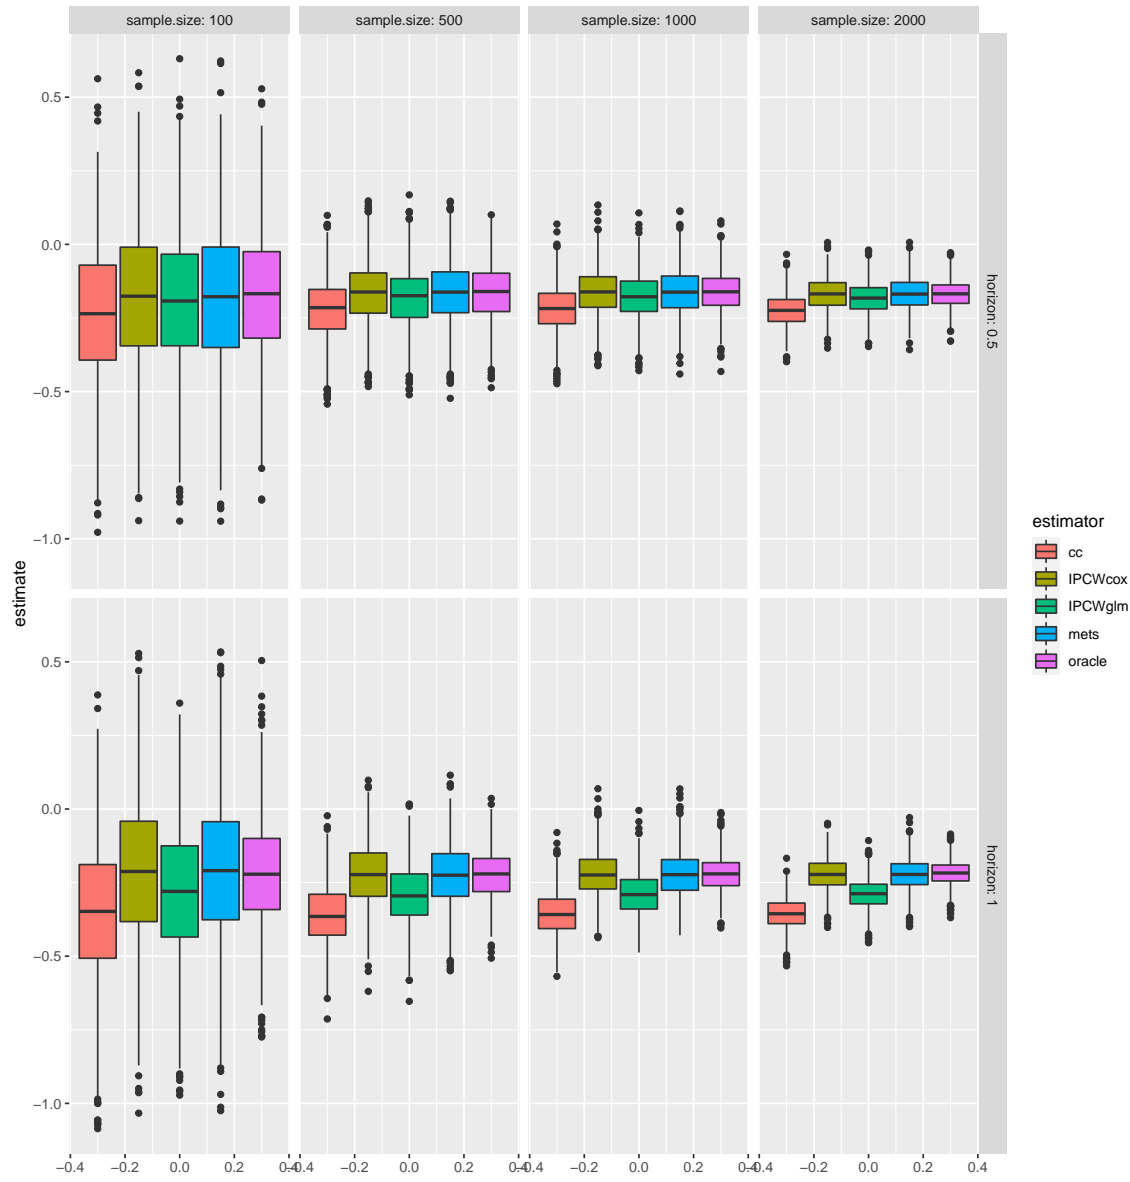


Figure 2: Comparison between the empirical distributions of the estimators (binary case) across sample size. Based on 1000 replicates.

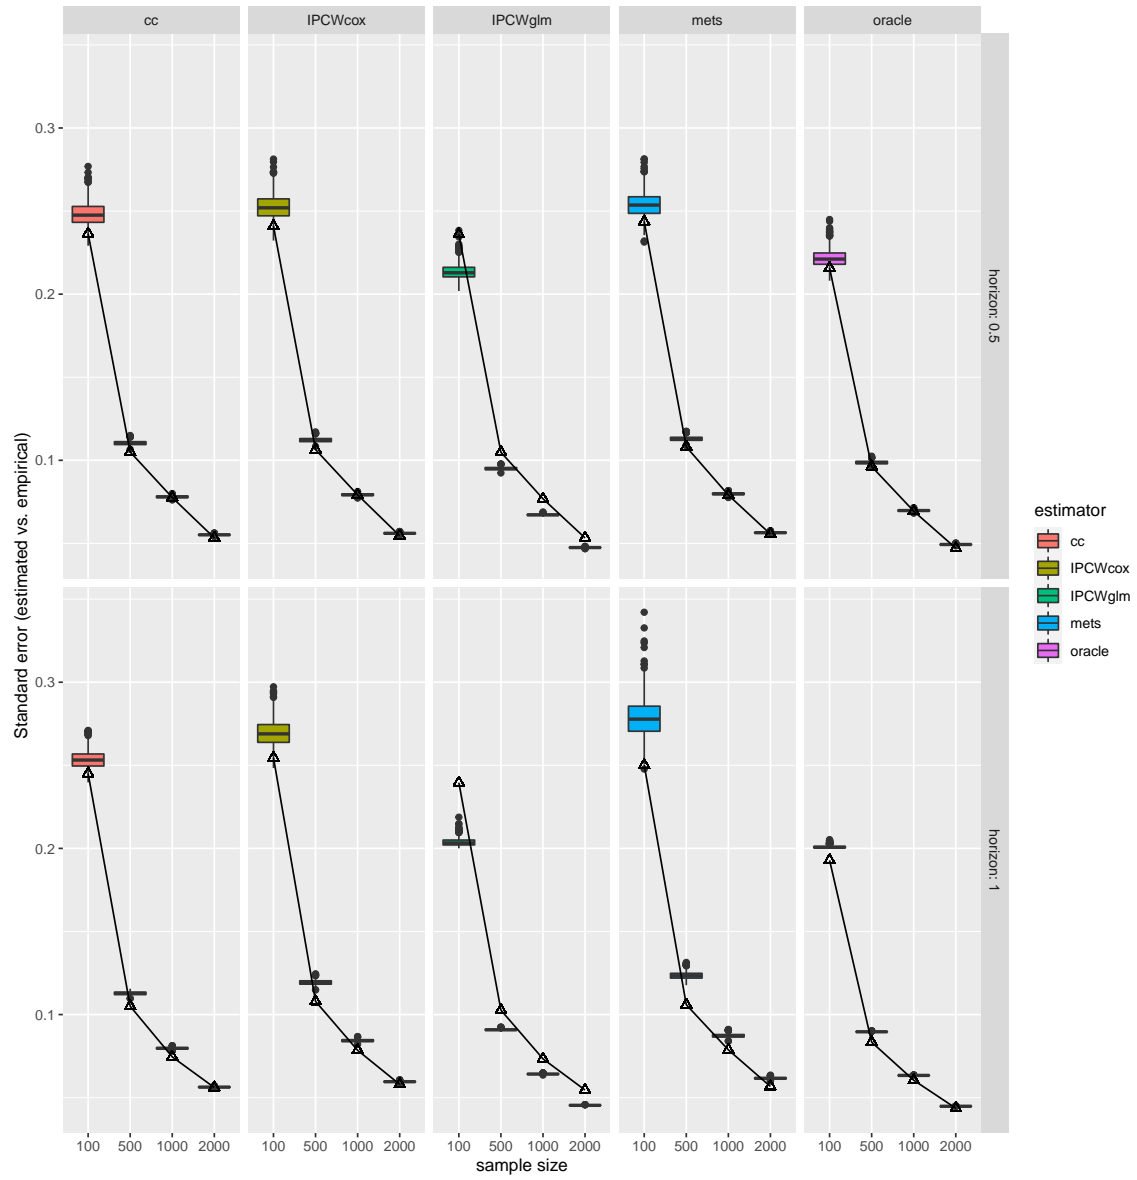


Figure 3: Comparison between the modeled standard errors of the estimates (boxplot) and the empirical ones (triangles linked by a line) across sample size. Based on 1000 replicates.